

Per-network single-step inference latency (CPU, 1 thread; bar=mean, whisker=p99)

